Problem Set 4

Stats 506, Fall 2018

Due: Monday December 10, 5pm

Instructions

- Submit the assignment by the due date via canvas. There is a maximum of 1 late day for this assignment.

- Use Rmarkdown to create and submit a single html or pdf with your answers to question 1-2 along with supporting evidence in the form of tables and graphs.

- All tables and graphs should be neatly labeled and appear polished.

- Question 1 and 2 ask you to use *R*. You should submit your code for each problem as `ps4_q1.R` and `ps4_q2_X.R`.

- You should submit a single compressed archive (`.zip`) which contains the following files:
    - `ps4.pdf` or `ps4.html`
    - `ps4.Rmd`
    - `ps4_q1.R`
    - `ps4_q2_funcs.R`, `ps4_q2a.R`, `ps4_q2b.R`, `ps4_q2c.R`
    - `run_ps4_q2b.pbs`, `run_ps4_q2c.pbs`
    - `ps4_q2b.Rout`, `ps4_q2c-X.Rout` (X = 1, 2, 4).
    - `ps4_q3.sas`, `ps4_q3c.csv`, `ps4_q3d.csv`
  All files should be executable without errors.

- All files read, sourced, or referred to within scripts should be assumed to be in the same working directory (`./`).

- Your code should be clearly written and it should be possible to assess it by reading it. Use appropriate variable names and comments. Your style will be graded using the style rubric (./StyleRubric.html) [15 points].

- Some of these exercises may require you to use commands or techniques that were not covered in class or in the course notes. You can use the web as needed to identify appropriate approaches. Part of the purpose of these exercises is for you to learn to be resourceful and self sufficient. Questions are welcome at all times, but please make an attempt to locate relevant information yourself first.

Question 1 [20 points]

Use the Lahman baseball data previously seen in the SQL notes
(https://jbhender.github.io/Stats506/F18/Intro_to_SQL.html) to answer this question. Your answer should be a
single SQL query, but may require anonymous tables created using nested queries.

Write an SQL query to construct a table showing the *all-time* leader in *hits* ("H" from the "batting" table) for
each birth country ("birthCountry" in the "master" table). An *all-time* leader is the player ("playerID") with the
most total hits across all rows (e.g. seasons/stints). Limit your table to players/countries with at least 200 hits
and order the table by descending number of hits. Create a nicely formatted table with the following columns
as your final output: Player (nameFirst nameLast), Debut (debut), Country of Birth (birthCountry), Hits (H).

Question 2 [40 points]

In this question you will modify your answer to Problem Set 3, Question 2 (PS3 Q2) to practice parallel,
asynchronous, and batch computing. Copy the functions from part a and c of PS3 Q2 to a new file
`ps4_q2_funcs.R`

In each of the parts below, let $\beta \in \mathbb{R}^{100}$ be defined so that

$$\beta_i = \begin{cases} .1, & i \le 10, \\ 0, & \text{else.} \end{cases}$$

and Σ be block diagonal with $\Sigma_{ij} = \rho\beta_i\beta_j$ when $i \ne j$ and $\Sigma_{ii} = 1$. (You may also use β as in PS3 Q2 and
rescale in any other way that results in a positive definite Σ.)

Create a table or plot for your results from each part.

a. Write an R script `ps4_q2a.R` that sources `ps4_q2_funcs.R`, and then uses `mclapply` to run parallel
   simulations for $\rho \in \{.25i\}_{i=-3}^{3}$.
   Let $\sigma = 1$ and use 10,000 Monte Carlo replications. Reorganize the results into a long data frame
   `results_q4a` with columns: "rho", "sigma", "metric", "method", "est", and "se". "Metric" should
   contain the assessment measure: FWER, FDR, Sensitivity, or Specificity and "method" the multiple
   comparison method used. The columns "est" and "se" should contain the Monte Carlo estimate and its
   standard error, respectively.

b. Use your script from part a as the basis for a new script `ps4_q2b.R`. Setup a 4 core cluster using
   `doParallel` and then use nested foreach loops to run simulations for $\rho \in \{.25i\}_{i=-3}^{3}$ and
   $\sigma = \{.25, .5, 1\}$. Reshape the results as before into `results_q4b` saved to a file
   `results_q4b.RData`. Use a PBS file to run this script on the Flux cluster.

c. Modify your script from part a to create `ps4_q2c.R` which reads the following arguments from the
   command line: `sigma`, `mc_rep`, and `n_cores`. Also modify the script to use the `futures` package
   for parallelism. Use a PBS file to run this script as a job array for $\sigma = \{.25, .5, 1\}$. *Hint: see the answer
   at this page (https://stackoverflow.com/questions/12722095/how-do-i-use-floating-point-division-in-
   bash) for how to convert* `$PBS_ARRAYID` *to* `sigma`.

Question 3 [25 points]

For this question you should use the 2016 Medicare Provider Utilization and Payment data available here here (https://www.cms.gov/Research-Statistics-Data-and-Systems/Statistics-Trends-and-Reports/Medicare-Provider-Charge-Data/Physician-and-Other-Supplier.html).

a. Put the data into a folder `./data` and then follow the instructions to read this data into SAS.

b. Use one or more data steps to reduce the data set to those rows with "MRI" in the 'hcpcs_description' field and where 'hcpcs_code' starts with a 7.

c. Use proc means or proc summary (as needed) to determine the MRI procedures with the highest volume, highest total payment, and highest average payment among the procedures represented here.

d. Repeat part b-c using PROC SQL.

e. Export the results from "c" and "d" to csv and verify that they match. You do *not* need to produce a nice table within your solution document.